

Artasam Bin Rashid

AI/ML Engineer · LLM Agents · MLOps · Generative AI

artasambinrashid@gmail.com | linkedin.com/in/artasam-bin-rashid | github.com/Artasam | Portfolio | Islamabad, Pakistan

SUMMARY

AI/ML Engineer specializing in LLM-powered agentic systems, RAG pipelines, and end-to-end MLOps. Built production multi-agent architectures with LangGraph (7-node StateGraph) and automated full ML lifecycles—data ingestion through CI/CD retraining—using GitHub Actions and Hopsworks Feature Store. Fine-tuned Stable Diffusion Inpainting for domain-specific generative vision. Seeking a role to apply deep expertise in LLM orchestration, autonomous agents, and scalable model deployment.

TECHNICAL SKILLS

Languages: Python, SQL

ML / DL Frameworks: PyTorch, TensorFlow, Scikit-Learn, Hugging Face Transformers, OpenCV, SHAP

LLM & Agents: LangChain, LangGraph, RAG, Groq API, OpenAI API, Hugging Face S-BERT, Prompt Engineering, LLM Evals

Generative AI: Stable Diffusion (fine-tuning / inpainting), Diffusion Models, Sentence Transformers

MLOps & Infra: MLflow, GitHub Actions (CI/CD), Docker, Hopsworks Feature Store, Streamlit, FastAPI

Vector & Search: Chroma, cosine similarity, TF-IDF, semantic search

Cloud & Tools: AWS (AI Practitioner), Git, spaCy, LightGBM, XGBoost

PROJECTS

Search Engine Agent — Multi-Source AI Research Assistant with YouTube Learning Mode [GitHub]

- Engineered a LangGraph ReAct-loop agent querying DuckDuckGo, Wikipedia, and Arxiv with dynamic switching across 4 Groq LLMs (Llama 3.3 70B / Mixtral) and per-model TPM budget management, reducing average research session time by ~60% vs. manual search.
- Built YouTube Learning Mode with 5-strategy transcript extraction, AI-generated summaries, and multi-turn video Q&A with conversation memory; session analytics and JSON/Markdown export.

LinkedIn Post Automation Agent — Production Multi-Agent CI/CD Pipeline [GitHub]

- Architected a 7-node LangGraph StateGraph pipeline autonomously generating and publishing daily LinkedIn content via GitHub Actions; **100% uptime over 60+ days**, zero manual intervention.
- 6-source news aggregator, 38-keyword relevance scoring, anti-repetition deduplication, 2-tier safety pipeline, and full LinkedIn REST API v2 integration.

AI Virtual Try-On System (Final Year Project)

- Fine-tuned Stable Diffusion Inpainting (stable-diffusion-v1-5) with OpenPose, body segmentation, and geometric garment warping; photorealistic overlays across 12+ body poses, **sub-3s on T4 GPU**.

Air Quality Index Prediction System (Rawalpindi) — Serverless MLOps Pipeline [GitHub]

- XGBoost achieved **RMSE 0.65, R² 0.99** on 72-hour AQI forecasting, outperforming Random Forest by 20%; SHAP explainability and real-time hazard alerts included.
- Automated best-model registration to Hopsworks with hourly feature ingestion and daily retraining via GitHub Actions CI/CD — zero manual intervention.

WORK EXPERIENCE

Data Science Intern · 10Pearls Pakistan (Remote) *Dec 2025 – Feb 2026*

- Architected production AQI forecasting system on Open-Meteo real-time API; XGBoost achieved **RMSE 0.66, R² 0.99**, outperforming Random Forest by 20%.
- Integrated Hopsworks Feature Store for feature versioning and model registry; GitHub Actions CI/CD eliminated manual retraining overhead across **60+ consecutive operational days**.
- Engineered time-based, rolling-statistics, and lag features across a fully serverless stack with hourly ingestion and daily retraining pipelines.

AI/ML Intern · DAM TECHHUB (Remote) *Jul 2025 – Sep 2025*

- Built AI-powered Resume Analyser: Logistic Regression + TF-IDF classifier achieving **87% accuracy** across 24 job categories on a 2,400-resume dataset; Streamlit deployment with sub-1s inference latency.
- Integrated Hugging Face S-BERT for cosine semantic similarity matching — **23% precision improvement** over keyword-only baseline.
- Extended platform with ATS simulation across 6 systems (Taleo, Workday, Greenhouse, iCIMS, Lever, Generic) and Groq Llama 3.3 70B AI bullet-polish rewriter.

EDUCATION

BSc Artificial Intelligence · National University of Modern Languages (NUML), Islamabad *Feb 2022 – Jan 2026*

- Final Year Project:** AI Virtual Try-On System — Stable Diffusion Inpainting fine-tuning with pose estimation and geometric garment warping
- Relevant Coursework:** Deep Learning, NLP, Computer Vision, Probabilistic ML, Reinforcement Learning, Data Structures & Algorithms

CERTIFICATIONS & VIRTUAL EXPERIENCE

- LangChain for LLM Application Development** — DeepLearning.ai (Andrew Ng), 2025
- AWS AI Practitioner** — Amazon Web Services, 2025
- GenAI Data Analytics Virtual Experience** — Tata Group via Forage, 2025–2026 — Applied GenAI to financial risk workflows: EDA, ML-based delinquency prediction, customer segmentation, and AI-powered collections strategy design.