

ARTASAM BIN RASHID

AI Engineer — LLM Agents & Automation — Prompt Engineering — RAG — Generative AI

artasambinrashid@gmail.com | linkedin.com/in/artasam-bin-rashid | github.com/Artasam | Portfolio | RWP / ISB, Pakistan

SUMMARY

AI/ML Engineer specializing in prompt engineering, LLM agent orchestration, and Retrieval-Augmented Generation (RAG), with applied experience building autonomous workflows that automate manual reporting and data-processing tasks. Built production multi-agent systems with LangGraph (7-node StateGraph) and end-to-end MLOps pipelines (data ingestion through CI/CD retraining) using GitHub Actions and Hopsworks Feature Store. Applied GenAI to financial-risk workflows (EDA, delinquency prediction, customer segmentation) in a Tata Group data-analytics engagement. Fine-tuned Stable Diffusion Inpainting for domain-specific generative vision. Seeking an AI Engineer role applying LLM orchestration, prompt engineering, and automation to real-world financial and enterprise data workflows.

PROFESSIONAL EXPERIENCE

Machine Learning Intern

Jul 2026 – Sep 2026

FlyRank AI (Remote)

- Queried and processed multi-tenant search performance data from a ~79M-row distributed warehouse (DuckDB/SQL), auditing 92,548 content pages across 40 enterprise clients with zero local storage bottlenecks.
- Designed and deployed a domain-informed ranking baseline (striking-distance position weights, non-linear CTR shields) achieving 0.64 Precision@50 (1.73x lift over 28.6% base rate) and 0.75 Precision@20 on unseen enterprise domains.
- Identified and proved a critical 50-percentage-point memorization trap (Random Split P@20: 0.90 vs. Grouped Split: 0.40), preventing production deployment of an overfitted 4-feature LightGBM model on new client domains.
- Architected a zero-leakage temporal evaluation harness (15-day historical feature window vs. 16-day future label window, 0-day overlap), verified against 7 internal decision flags.
- Shipped an automated content-action playbook with interpretable reason codes and a 5-step human-in-the-loop review framework to guide editorial refresh operations for high-ROI search queries.

Data Science Intern

Dec 2025 – Feb 2026

10Pearls Pakistan (Remote)

- Architected a production forecasting system automating a manual reporting process end-to-end on a real-time API feed; XGBoost model achieved RMSE 0.66, R^2 0.99, outperforming Random Forest by 20%.
- Integrated Hopsworks Feature Store for feature versioning and model registry; GitHub Actions CI/CD eliminated manual retraining overhead across 60+ consecutive operational days.
- Engineered time-based, rolling-statistics, and lag features across a fully serverless data-processing stack with hourly ingestion and daily automated retraining pipelines.

AI/ML Intern

Jul 2025 – Sep 2025

DAM TECHHUB (Remote)

- Built an AI-powered Resume Analyser for automated text extraction and classification: Logistic Regression + TF-IDF classifier achieving 87% accuracy across 24 job categories on a 2,400-document dataset; deployed via Streamlit with sub-1s inference latency.
- Applied prompt engineering with Groq GPT-OSS-120B to build an AI bullet-polish rewriter, and integrated Hugging Face S-BERT for semantic similarity matching — a 23% precision improvement over keyword-only extraction.
- Extended the platform with ATS-simulation compatibility across 6 systems (Taleo, Workday, Greenhouse, iCIMS, Lever, Generic), validating structured-data extraction accuracy.

GenAI Data Analytics Virtual Experience

July 2025 – April 2026

Tata Group (via Forage)

- Applied Generative AI to financial risk workflows: exploratory data analysis, ML-based delinquency prediction, customer segmentation, and AI-powered collections-strategy design.

TECHNICAL SKILLS

Prompt Engineering & LLM Orchestration: Prompt Engineering, LLM Evals, LangChain, LangGraph, RAG, OpenAI API, Groq API, Hugging Face Transformers, S-BERT

Process Automation & Agents: Autonomous LLM agent workflows, multi-agent StateGraph design, GitHub Actions CI/CD, scheduled/serverless pipelines, ActivePieces (workflow automation — triggers, branching, API steps)

Data Processing & Analytics: Python, SQL, data cleaning & feature engineering, SHAP explainability, semantic search, TF-IDF, cosine similarity

Model Customization: RAG pipeline design, Stable Diffusion fine-tuning, feature engineering with Hopsworks Feature Store, MLflow

Vector & Search: Chroma (vector DB), semantic search, cosine similarity, TF-IDF

Deployment & MLOps: Docker, FastAPI, Streamlit, MLflow, GitHub Actions, AWS (AI Practitioner)

ML / DL Frameworks: PyTorch, TensorFlow, Scikit-Learn, Hugging Face Transformers, OpenCV, spaCy, LightGBM, XGBoost

PROJECTS

LinkedIn Post Automation Agent | Production Multi-Agent CI/CD Pipeline | [GitHub]

- Architected a 7-node LangGraph StateGraph pipeline as an autonomous LLM workflow, replacing manual content publishing via GitHub Actions; 100% uptime over 60+ days with zero manual intervention.
- Built a 6-source data aggregator with 38-keyword relevance scoring, anti-repetition deduplication, a 2-tier safety pipeline, and full LinkedIn REST API v2 integration.

Search Engine Agent | Multi-Source AI Research Assistant | [GitHub]

- Engineered a LangGraph ReAct-loop agent for autonomous multi-source research (DuckDuckGo, Wikipedia, Arxiv) with dynamic switching across 4 Groq LLMs and per-model TPM budget management, reducing manual research time by ~60%.
- Built a YouTube Learning Mode with 5-strategy transcript extraction, AI-generated summaries, multi-turn Q&A with conversation memory, and session analytics with JSON/Markdown export.

AI Virtual Try-On System | Final Year Project

- Fine-tuned Stable Diffusion Inpainting (stable-diffusion-v1-5) with OpenPose, body segmentation, and geometric garment warping; photorealistic overlays across 12+ body poses at sub-3s inference on a T4 GPU.

Air Quality Index Prediction System (Rawalpindi) | Serverless MLOps Pipeline | [GitHub]

- XGBoost achieved RMSE 0.65, R² 0.99 on 72-hour forecasting, outperforming Random Forest by 20%; included SHAP explainability visualizations and real-time hazard alerts.
- Automated best-model registration to Hopsworks with hourly feature ingestion and daily retraining via GitHub Actions CI/CD — zero manual intervention.

AI Hiring Pipeline Automation Agent — ActivePieces Workflow

- Building a scheduled, trigger-based automation pipeline in ActivePieces that fetches open roles from a careers page, deduplicates against stored state, and routes new job descriptions through email confirmation before automated resume scoring — covering the same trigger/branch/API-integration logic used in n8n and Make.

EDUCATION

National University of Modern Languages (NUML), IslamabadFeb 2022 – Jan 2026BS Artificial Intelligence

- Final Year Project: AI Virtual Try-On System — Stable Diffusion Inpainting fine-tuning with pose estimation and geometric garment warping.
- Relevant Coursework: Deep Learning, NLP, Computer Vision, Probabilistic ML, Reinforcement Learning, Data Structures & Algorithms.

CERTIFICATIONS

LangChain for LLM Application Development — DeepLearning.AI (Andrew Ng)2025

AWS AI Practitioner — Amazon Web Services (Udacity)2026